

Training Through the Inferno: An Inspectable OSRS Boss-Fight Benchmark for RL

Anonymous Authors

Abstract

Modern game RL benchmarks need fast training, inspectable failures, and enough simulator fidelity that policy mistakes mean something. We present an Old School RuneScape Inferno benchmark built in PufferLib 4. The Inferno is a 69-wave solo boss encounter with prayer switching, line-of-sight positioning, supplies, target priority, and a moving-shield final phase. These mechanics create long episodes, partial observability, sparse success, and failure modes that skilled players can read.

The contribution is a benchmark design pattern: player-readable state, structured actions, strict masks, replayable failures, behavior-level logs, and high-throughput recurrent training. The current environment exposes 744 symbolic observation features and a 9-head action space with 89 total discrete choices. The observation includes one mask entry per choice. We use PufferLib 4’s native recurrent RL stack to make large-agent, long-episode iteration practical. A compact recurrent checkpoint from the preceding compact Redemption action surface trained for 171.7M environment steps and reached 0.49 logged development win rate and 0.736 internal training score in the stored run summary. This run is development telemetry from an earlier compact action mapping, not a frozen-schema evaluation of the current 744-feature, 9-head, 89-choice surface. We use it only to show that this benchmark line can produce logged full clears and late-fight diagnostics.

1 Introduction and Contribution

The Old School RuneScape Inferno is a solo player-versus-monster challenge released in 2017. The player must survive 69 waves and defeat TzKal-Zuk to earn the infernal cape. The OSRS Wiki describes the activity as a no-restock encounter with three pillars that block projectiles during the waves and a mandatory shield mechanic in the final fight [1]. That readability makes the Inferno a useful game-RL benchmark. If the agent dies to a Jal-Zek, it may have missed a magic prayer. If a pillar falls, it likely ignored nibblers. If Zuk shield health drops, it failed movement, target priority, or healer timing.

This paper is about benchmark construction. We do not wrap a game binary and report a leaderboard number. We build a simulator whose state, actions, masks, renderer, replay path, and logs form one contract. The same game facts used for training must be available when explaining watched behavior. If a policy looks wrong in replay, the logs and simulator should explain why.

The benchmark fits RLVG along three axes. First, it is a game-derived task with long-horizon resource management and timing. Second, it supports behavioral evaluation: the benchmark exposes whether policy behavior matches player-readable task structure when scalar reward rises. Third, it depends on systems work. The environment became testable only when it could train, render, and replay fast enough to make debugging empirical.

2 Benchmark Surface

The Inferno environment uses symbolic observations instead of pixels. The current surface has 744 base observation features and a 9-head action space with 89 total discrete choices. The observation includes one mask entry per choice. One environment step is one OSRS game tick, or 0.6 seconds. Episodes can run for thousands of ticks. The native environment writes fixed-size observations and masks into PufferLib rollout buffers for thousands of parallel agents.

The observation vector contains player health, prayer, supplies, wave phase, gear state, combat stats, target state, weapon range, NPC slots, pillar health, pending hits, pending Zuk healer area attacks, and short movement forecasts. Positions are normalized and mostly egocentric or arena-relative. Categories use one-hot structure or slot membership. The design rule is simple: expose player-relevant state, then log whether the policy acts on it.

The action space is a structured interface over OSRS intent: movement, overhead prayer, target selection, gear switching, food, potion, spell, special attack, and offensive prayer. The target head points into observation slots instead of raw entity ids. The mask writer validates every action head and records zero-valid-head failures. Invalid

Component	Design choice	Why it matters
Observation	744 symbolic features	Exposes player-relevant state
Action	9 structured heads	Matches game intent
Masks	One entry per choice	Removes invalid client actions
Debug	Renderer, replay, logs	Makes failures inspectable
Training	4096 parallel agents	Keeps long episodes cheap
Evaluation	Normal-start summary	Separates full-run progress

Table 1: The benchmark surface is designed so training, replay, masks, and logs refer to the same game facts.




Figure 1: An Inferno wave 69 replay frame. The same contract elements used for training and debugging are visible: tick, target label, arena state, player and Zuk positions, HP and prayer counters, minimap, inventory, and action tabs.

actions are interface noise. The model should decide whether to run west, pray magic, tag a healer, cast barrage, or drink a restore. It should spend capacity on game decisions, not on learning that an absent NPC cannot be attacked. Reward and score are separate. Reward can be shaped, swept, and replaced. Score is intended to survive those edits. In this environment, reward paid for credit assignment through damage, healer tags, shield preservation, late-wave pressure, and final-phase priorities. The logged score in Table 3 is also separate from the archive `progress_score` used for Zuk-state selection. For full-start runs, it is win rate plus half-weighted wave progress on losses: $\text{winrate} + (1 - \text{winrate}) \cdot \text{mean_wave}/69 \cdot 0.5$. This explains why a 0.489635 win rate and 66.720871 mean wave produce a 0.736389 score.

3 Training and Iteration Loop

PufferLib was introduced as a way to make RL libraries and environments interoperate [7]. PufferLib 4 made this loop practical because the bottleneck is the full edit-train-watch cycle. The public Puffer docs describe a native backend with CUDA C kernels, static memory, CUDA graphs, async environment workers, fused operations, action masks, PufferNet policies, and Protein sweeps [2]. Joseph Suarez’s Puffer 4 engineering notes frame the same systems problem: small-model RL can spend too much time in framework overhead unless memory layout, kernel launch

overhead, and fused computation are designed together [3].

The recurrent model is PufferNet, which combines MinGRU with highway-style output mixing. MinGRU supports parallel-scan training across sequences and cheap recurrent inference during rollouts [4, 5]. This fits game tasks where the policy needs memory during rollout, but training must still use long enough sequences to learn delayed consequences.

The Inferno config uses 4096 parallel agents, two rollout buffers, 16 worker threads, a two-layer 512 hidden PufferNet, a 16 tick horizon, priority replay, and a replay ratio of 4. The horizon covers 9.6 seconds of game time. Protein sweeps search over learning rate, horizon, discounting, advantage settings, replay, vector size, model width, curriculum fractions, and reward weights. Protein models both score and wall-clock cost, which matters because a high score from a slow or unstable setting may be a poor development tool [6].

The hardest engineering lesson was that the simulator and learner could not be debugged separately. Training exposed missing observations, stale masks, and reward traps. Renderer checks exposed bugs in projectile anchoring, Zuk pathing, healer tags, and action timing. Human controls exposed when the action interface was not playable. Some fixes crossed every surface: prayer flick semantics touched environment logic, GUI, human input, rendering, tests, and trainer bindings, while Zuk healer work produced diagnostics before reward changes. The useful loop became: train, watch, find a player-readable failure, trace the simulator, add a contract test, and sweep again.

Observed failure	Visible signal	Contract added
Wrong prayer timing	Prayer correctness and death context	Pending-hit observations
Stale target options	Zero-valid-head counters	Strict target masks
Missed Zuk healers	Healer tag and kill logs	Healer phase diagnostics
Broken Zuk movement	Replay and human click pathing	Shield-path checks
Wrong projectile source	Rendered flight path	Anchored projectile endpoints

Table 2: Examples where policy watching, metrics, and simulator contracts found different classes of defect.

4 Checkpoint Evidence

The checkpoint result is historical development telemetry from the active branch, not a frozen-schema benchmark evaluation. It used a compact Redemption action mapping, state curriculum, priority replay, 4096 parallel agents, a two-layer 512 hidden PufferNet, and 171.7M environment steps. The checkpoint README records a 5-action PvE overhead head in which the compact mapping represented Redemption without the current explicit Redemption logit. Its action mapping differs from the current benchmark surface summarized in Table 1. The stored run summary reports the development telemetry in Table 3. The summary marks the final metrics as normal-start evaluation, but this artifact is downsampled, so we avoid treating its `env/n` field as a precise evaluation denominator here.

Metric	Value
Environment steps	171.7M
Logged win rate	0.489635
Internal training score	0.736389
Average final wave	66.720871
Average minimum Zuk HP, wins as 0	240.535599
Prayer correctness	0.853395
Fraction reaching Zuk healers	0.847019
Fraction killing all Zuk healers	0.718624
Final logged throughput	about 256k steps/s

Table 3: Stored development telemetry for the compact recurrent Inferno checkpoint.

Read these numbers as development evidence. The environment, action surface, and visual contract were still changing during development. The result shows logged full clears and frequent late-Zuk states in this benchmark line. It does not prove that the task is solved, that the checkpoint is robust across future environment versions, or that the policy has no simulator-specific habits.

This distinction is central to the benchmark. In an inspectable game task, success cannot be only a scalar. A policy that clears while wasting supplies, mistiming healers, or leaning on a simulator bug has not solved the benchmark in

the player-facing sense. The benchmark must keep these failures visible.

A stricter public benchmark release needs more than extra episodes. It should freeze the observation and action schema, pin the checkpoint metadata, run normal-start no-render evaluation with an explicit episode count, replay sampled wins and losses, and report behavioral counters alongside score. A policy should be judged on full clears, late-phase progress, and whether watched failures remain explainable from logged state.

5 Lessons and Limitations

First, game benchmarks need a player-facing contract. If policy behavior looks wrong in replay, the logs and simulator should explain why. This is not a cosmetic requirement. In the Inferno, rendering and replay caught bugs that scalar training curves would have hidden.

Second, strict action masks are part of the task definition. OSRS has many impossible actions: attacking missing targets, drinking empty potions, casting unavailable spells, switching to already equipped gear, or stepping into blocked tiles. Masking these actions does not make the game easy. It removes interface noise so the model can spend capacity on game decisions.

Third, reward is not score. Reward is an engineering tool for credit assignment. Score is the stable development yardstick. Keeping them separate made the project less brittle while reward weights and curricula changed.

Fourth, higher throughput made the development loop more empirical. PufferLib 4’s static memory, fused kernels, CUDA graphs, MinGRU, and sweep tooling were not aesthetic choices. They made it cheap enough to find out whether a simulator fix or observation change actually helped [3, 4, 6].

The main limitation is fidelity. This is not a full OSRS client. It is a simulator for the encounter mechanics that matter for training, plus cache-derived assets for visual inspection. Expert knowledge makes failures interpretable, but it can also hide assumptions in the simulator. The safest path is to keep tests, render checks, replay checks, and human play close to the training surface.

The OSRS Inferno environment suggests a benchmark pattern for long-horizon game RL: player-readable state, structured actions, strict masks, replayable failures, and high-throughput recurrent training. A benchmark should report whether an agent wins and make the agent’s mistakes legible.

References

- [1] OSRS Wiki. Inferno. <https://oldschool.runescape.wiki/w/Inferno>
- [2] PufferAI. PufferLib docs. <https://puffer.ai/docs.html>
- [3] Joseph Suarez. The Engineering behind PufferLib 4.0’s 20,000,000 step/second RL Training. <https://x.com/jsuarez/article/204150187922284776>
- [4] Joseph Suarez. A High Throughput Recurrent Network for PufferLib 4.0. <https://x.com/jsuarez/article/2042273938592407674>
- [5] Leo Feng, Frederick Tung, Mohamed Osama Ahmed, Yoshua Bengio, and Hossein Hajimirsadeghi. Were RNNs All We Needed? arXiv:2410.01201, 2024. <https://arxiv.org/abs/2410.01201>
- [6] Joseph Suarez. Visualizing 20k RL Experiments with Constellation. <https://x.com/jsuarez/article/2041896392587722885>
- [7] Joseph Suarez. PufferLib: Making Reinforcement Learning Libraries and Environments Play Nice. arXiv:2406.12905, 2024. <https://arxiv.org/abs/2406.12905>